Granite-3.3:8B — Code Completion Evaluation

Comparison of Hole Filler Template vs. Fill In the Middle Template

Appendix: Prompt Design (for Reference)

• Hole Filler Template:

Prefix only; cursor at end of code block.

Prompts used: https://github.com/IBM-OSS-Support/Continue.dev-Granite-manual-test-cases/blob/main/tab-autocomplete/granite3.3-8b/usecases/prefix-only/prefix-only-usecases.py

• FIM Template:

Prefix (before hole), Suffix (after hole), cursor at the gap.

Prompts used: https://github.com/IBM-OSS-Support/Continue.dev-Granite-manual-test-cases/blob/main/tab-autocomplete/granite3.3-8b/usecases/fim/fim-usecases.py

Executive Summary

• Hole Filler Template (HFT):

- Strengths: Reliable for sequential code completions, strong with Python syntax and functional code, robust for simple to moderately complex code.
- Weaknesses: Prone to hallucinations in structured/multi-branch logic, verbose in certain scenarios.

Fill In the Middle (FIM):

- Strengths: Excels at filling isolated code holes (e.g., lambdas, recursion, exception logic), ideal for small focused completions, integrates well in refactoring workflows.
- Weaknesses: Struggles with complex, multi-branch logic and structural completions (e.g., tax slabs, pandas aggregations, OOP methods); sometimes produces incomplete or vague output.

Evaluation Summary Table

Test Case	Hole Filler Verdict	FIM Verdict
1. Nested Conditions (Tax Slabs)	A Partial: Syntax ok, logic error in slab calc	X Broken: Dead code, misses slab, syntax err
2. Lambda + Filter	▼ Excellent: Clean, correct, idiomatic	Excellent: Clean, correct, idiomatic

3. Pandas Chaining	✓ Good: Uses NamedAgg, minor format issue	X Vague: "all columns" non- Python output
4. Exception Handling with Custom Message	Solid: Logical, correct, verbose	✓ Excellent: Clean, concise, correct
5. Class with Dunder/Bonus	Good: Dunder and bonus logic are correct	X Broken: Incomplete, missing return, logic
6. Recursive Function (Factorial)	Excellent: Canonical recursion, clean	Excellent: Canonical recursion, clean

Detailed Comparison by Use Case

Nested Conditions (Tax Slabs)

Hole Filler:

Partial logic, hardcodes slab values, does not fully respect business logic, but structure is valid Python.

FIM:

Misses a whole branch, inserts dead code, and outputs syntactically invalid Python.

Lambda + Filter

• Hole Filler:

Generates a correct, idiomatic lambda filter for even numbers.

• FIM:

Also generates correct lambda. Both modes perform equally well.

Pandas Chaining

• Hole Filler:

Uses advanced API (NamedAgg), correct overall structure, only minor formatting issues.

• FIM:

Produces non-Python output ("all columns"), does not provide a real aggregation statement.

Exception Handling with Custom Messages

Hole Filler:

Robust, verbose, but logic is correct.

• FIM:

Clean, concise, and correct — in fact, more succinct than HFT.

Class Definitions with Dunder/Bonus

Hole Filler:

- Generates correct __init__ and __str__.
- Implements calculate_bonus logic as intended (salary check, bonus calculation, returns correct value).
- Minor verbosity at end, but not incorrect.
- **Result: Solid performance** for both OOP syntax and business logic.

• FIM:

- Output is incomplete, missing both condition and return.
- Result: Fails to provide usable code for this scenario.

Recursive Logic

Hole Filler:

Outputs canonical recursive implementation; perfect.

• FIM:

Same as HFT; perfect recursion.

Strengths & Weaknesses

Hole Filler Template (HFT)

Strengths:

- Good at classic code completion and generating boilerplate or sequential code.
- Handles simple functions, recursion, and basic functional programming patterns well.
- Can generate complex API calls if context is linear.
- Excels at both classic and OOP code completions, including dunder methods and conditional logic inside methods.

Weaknesses:

- Hallucinates or fumbles structured/branching logic.
- Tends to generate verbose code (especially with print statements).

Fill In the Middle (FIM) Template

Strengths:

- Excels at targeted, in-place code hole filling (e.g., filling in a lambda, a recursive call, or a single exception raise).
- More concise for isolated code blocks.

Integrates well in a refactoring workflow or patch-based code editing.

Weaknesses:

- Unreliable for multi-branch, multi-step, or structural completions (e.g., tax slab logic, pandas aggregations, full method bodies).
- Sometimes produces vague or invalid output when the required logic is complex.
- Suffers on OOP patterns and logic-heavy completions.

Overall Verdict

- For most real-world developer workflows:
 - Use Hole Filler Template for general code completion, larger blocks, or full function/method generation.
 - Use FIM Template for targeted patching, simple insertions, or automated refactoring tasks that require context before and after the hole.

Model Quality:

- Granite-3.3:8B is strong on simple/medium code completion, recursion, lambdas, and straightforward exception logic in both modes.
- It needs improvement in class structure, complex branching, and context-heavy or "abstract" tasks (like custom aggregation).

Recommendations

Use Case Fit:

- **Hole Filler:** For new function writing, classic autocomplete, or when the task is to finish a code block given the start.
- **FIM:** For editing/filling gaps in the middle of existing code, quick patching, or tight refactoring loops.